

BoostingBERT: Integrating Multi-Class Boosting into BERT for NLP Tasks

Tongwen Huang
Tencent Corp, China
towanhuang@tencent.com

Qingyun She
Sina Weibo Corp, China
qingyu@staff.weibo.com

Junlin Zhang
Sina Weibo Corp, China
junlin6@staff.weibo.com

Abstract

As a pre-trained Transformer model, BERT (Bidirectional Encoder Representations from Transformers) has achieved ground-breaking performance on multiple NLP tasks. On the other hand, Boosting is a popular ensemble learning technique which combines many base classifiers and has been demonstrated to yield better generalization performance in many machine learning tasks. Some works have indicated that ensemble of BERT can further improve the application performance. However, current ensemble approaches focus on bagging or stacking and there has not been much effort on exploring the boosting. In this work, we proposed a novel Boosting BERT model to integrate multi-class boosting into the BERT. Our proposed model uses the pre-trained Transformer as the base classifier to choose harder training sets to fine-tune and gains the benefits of both the pre-training language knowledge and boosting ensemble in NLP tasks. We evaluate the proposed model on the GLUE dataset and 3 popular Chinese NLU benchmarks. Experimental results demonstrate that our proposed model significantly outperforms BERT on all datasets and proves its effectiveness in many NLP tasks. Replacing the BERT base with RoBERTa as base classifier, BoostingBERT achieves new state-of-the-art results in several NLP Tasks. We also use knowledge distillation within the "teacher-student" framework to reduce the computational overhead and model storage of BoostingBERT while keeping its performance for practical application.

1 Introduction

Until recently, the NLP community is witnessing a dramatic paradigm shift toward the pre-trained deep language representation model, which achieves state of the art in various NLP tasks such as question answering (Zhang et al., 2019; Nogueira and Cho, 2019), sentiment classification (Sun et al., 2019)

and relation extraction (Han et al., 2019).

By pre-training on a large corpus on tasks such as language modeling, deep language representations are generated that can serve as a starting point to a variety of NLP tasks. Examples of pre-trained models includes ELMo (Peters et al., 2018), ULMFiT (Howard and Ruder, 2018), OpenAI GPT (Radford et al., 2018) and most recently BERT (Devlin et al., 2018). Among them, BERT outperforms its predecessors and it consists of two stages: first, BERT is pre-trained on vast amounts of text, with an unsupervised objective of masked language modeling and next-sentence prediction. Next, this pre-trained network is then fine-tuned on task-specific labeled data.

On the other hand, ensemble is a successful approach to reduce the variance of sub-models by training multiple sub-models to combine the predictions from them. Recently, some ensemble approaches using the pre-trained transformer as the base classifier are proposed (Fajcik et al., 2019; Liu et al., 2019a). However, these works focus on bagging or stacking and there has not been much effort on exploring the boosting.

In this work, we propose a BoostingBERT model to introduce the multi-class boosting into BERT. We all know that some training instances with specific characteristics are more difficult to classify in many NLP tasks. We can use boosting to add extra pre-trained base Transformer and let these pre-trained base classifiers pay more attention to these difficult instances for a better task performance. In this way, BoostingBERT model gains the benefits of both the pre-training language knowledge and boosting ensemble.

In summary, this paper has four major contributions:

1. We propose a novel model named BoostingBERT that integrates multi-class boosting into

BERT. As far as we know, this work is the first one to prove that **boosting can be used to enhance the performance of BERT, instead of bagging or stacking**. Our experimental results also demonstrate that BoostingBERT outperforms the bagging BERT constantly.

2. We compare two approaches making use of the base Transformer classifier in BoostingBERT model: weights privacy vs. weights sharing. Experiment results demonstrate that the former one constantly outperforms the latter one.
3. We conduct extensive experiments on the GLUE dataset and 3 popular Chinese NLU benchmarks. The experiment results show that BoostingBERT significantly outperforms BERT on all tasks and demonstrates its effectiveness in many NLP tasks. We also find that BoostingBERT is particularly useful for the tasks with little training data.
4. Considering BoostingBERT’s large number of parameters and long inference time, we use **knowledge distillation** within the "teacher-student" framework to reduce the computational overhead and model storage of BoostingBERT while keeping its performance.

The rest of this paper is organized as follows. Section 2 introduces some related works which are relevant with our proposed model. We introduce our proposed Boosting BERT model in detail in Section 3. The experimental setup and results on GLUE dataset and several Chinese benchmark datasets is presented in Section 4 and Section 5. Section 6 concludes our work in this paper.

2 Related Work

2.1 Pre-trained Models in Natural Language Processing

Inspired from the computer vision field, where ImageNet (Deng et al., 2009) is used to pre-train models for other tasks (Huh et al., 2016), many pre-trained general-purposed language encoders have generated a lot of interest in the NLP community. In contrast to fixed word embedding such as Word2Vec (Mikolov et al., 2013) or Glove (Pennington et al., 2014), the newer embedding generally incorporates both the left and right context.

Howard (Howard and Ruder, 2018) proposed a general transfer learning method, Universal Language Model Fine-tuning (ULMFiT), with the key

techniques for fine-tuning a language model. Peters (Peters et al., 2018) introduce embedding from Language Models (ELMo), an approach for learning high-quality, deep contextualized representations using bidirectional language models. They achieve large improvements on six different NLP tasks.

Radford (Radford et al., 2018) proposed that by generative pre-training of a language model on a diverse corpus of unlabeled text, large gains on a diverse range of tasks could be realized. BERT (Devlin et al., 2018) is the most recent inclusion to these models, where it uses a deep bidirectional transformer trained on masked language modeling and next sentence prediction objectives. It exceeds state of the art by a wide margin on multiple natural language understanding tasks.

2.2 Ensemble Learning

Ensemble learning uses multiple learning algorithms to obtain better predictive performance than could be obtained from any of the constituent learning algorithms alone by reducing the variance of single model. It is also one of the most popular approaches used by winners in many machine learning competitions. Bagging, boosting and stacking are three most widely used ensemble types.

Bagging, developed by Breiman (Breiman, 1996), is a machine-learning method that uses bootstrapping to achieve differences between models. Models are trained on different subsets of the training data naturally through the use of resampling methods such as cross-validation and the bootstrap, designed to estimate the average performance of the model generally on unseen data. The models used in this estimation process can be combined in what is referred to as a resampling-based ensemble, such as a cross-validation ensemble or a bagging ensemble.

Boosting (Freund and Schapire, 1997) assumes the availability of a “weak” or base learning algorithm which, given labeled training examples, produces a “weak” or base classifier. The goal of boosting is to improve the performance of the base learning algorithm. The key idea behind boosting is to choose training sets for the base classifier in such a fashion as to force it to infer something new about the data each time it is called. The fusion algorithm will finally combine many base classifiers into a single classifier whose prediction power is strong.

Stacking involves training an entirely new model to combine the predictions of several other sub-

models. First, all of the sub-models are trained using the available data, then a combiner algorithm is trained to make a final prediction using all the predictions of the sub-models as additional inputs.

2.3 Ensemble Models Based on BERT

Devlin (Devlin et al., 2018) reported that BERT shows significant increase in improvements on many NLP tasks, and subsequent studies have shown that BERT is also effective on harder tasks such as open-domain question answering (Zhang et al., 2019), multiple relation extraction (Han et al., 2019), and information retrieval (MacAvaney et al., 2019).

Recently, some ensemble approaches using the pre-trained transformer as the base classifier are also proposed. However, these ensemble focus on bagging or stacking. For example, Fajcik (Fajcik et al., 2019) introduced a bagging ensemble in determining the rumor stance with pre-trained deep bidirectional transformers. The base transformer classifiers differ just by learning rate and they propose 4 different fusion approaches to increase the performance. Liu (Liu et al., 2019a) proposed an ensemble framework to address the fake news classification challenge in ACM WSDM Cup 2019. In the proposed ensemble framework, they trained a three-level model to perform the fake news classification. The first level of the framework contains 25 BERT model with a blending ensemble strategy. Huang (Huang et al., 2019) proposed a Hierarchical LSTMs for Contextual Emotion Detection (HRLCE) model to classify the emotion of an utterance given its conversational context. HRLCE is also a stacking ensemble framework which combines the results generated by BERT and HRCLE to achieve better task performance compared to BERT.

3 Our Proposed Model

In this section, we describe our proposed Multi-class Boosting solution which is built on top of the pre-trained language encoder BERT in detail.

3.1 BERT

BERT is designed to learn deep bidirectional representations by jointly conditioning both the left and right contexts in all layers. The BERT model is pre-trained with two approaches on large-scale unlabeled text, i.e., masked language model and next sentence prediction. BERT proposes a “masked language model” (MLM) objective by masking some

percentage of the input tokens at random, and predicting the masked words based on its context. The pre-trained BERT model provides a powerful context-dependent sentence representation and can be used for various target tasks, through the fine-tuning procedure, similar to how it is used for other NLP tasks.

The model architecture of BERT is a multi-layer bidirectional Transformer encoder based on the original Transformer model (Vaswani et al., 2017). The input representation is a concatenation of Word-Piece embedding (Wu et al., 2016), positional embedding and the segment embedding. A special classification embedding ([CLS]) is inserted as the first token and a special token ([SEP]) is added as the final token. The BERT model can also encode multiple text segments simultaneously, allowing it to make judgments about text pairs. To adapt BERT for specific tasks, all parameters of BERT are fine-tuned jointly by predicting a task-specific label with the task-specific output layer to maximize the log-probability of the correct label.

3.2 Multi-Class Boosting Based on BERT

Boosting, e.g., AdaBoost, is a popular ensemble learning technique, which combines many “weak” or base classifiers and has demonstrated to yield better generalization performance in many applications. Boosting takes as input a training set $\mathcal{D} = \{(x_1, y_1), \dots, (x_i, y_i), \dots, (x_m, y_m)\}$, where each x_i belongs to some instance space X , and each label y_i is in some label set Y . AdaBoost calls a given “weak” or base learning algorithm repeatedly in a series of rounds $t = 1, \dots, i, \dots, T$. One of the main ideas of the boosting is to maintain a distribution or set of weights over the training set. The weight of this distribution on training example i on round t is denoted $\mathcal{D}_t(i)$. Initially, all weights are set equally, but on each round, the weights of incorrectly classified examples are increased so that the “weak” or base learner is forced to focus on the hard examples in the training set.

We intent to combine the Boosting with BERT in this work. Though original AdaBoost is a successful technique for solving the two-class classification problem, most NLP tasks belong to multi-class classification problem. So a multi-class version boosting algorithm is needed firstly for further work. On the other hand, the classic base classifier fusion strategy of AdaBoost is the weighted voting by base classifiers. We deem it's relatively simple



to effectively ensemble the complex base models such as BERT and we need a more complicated fusion approach. So we design BoostingBERT as a two stages process: The first stage introduces the multi-class classification ability into BERT to train the pre-trained Transformer classifiers in boosting way; In the second stage, we train a fusion network to ensemble each Transformer encoder after fixing the parameters of base classifiers. Because the performance of **fusion network** is comparable with or slightly better than the commonly used weighted voting strategy, so we use the strategy during this work.

3.2.1 First Stage: Base Classifier Training

In going from two-class to multi-class classification, most boosting algorithms have been restricted to reducing the multi-class classification problem to multiple two-class problems. Zhu (Zhu et al., 2009) proposed a new algorithm that directly extends the AdaBoost algorithm to the multi-class case without reducing it to multiple two-class problems. Surprisingly, this algorithm is almost identical to AdaBoost but with a simple modification. Inspired by Zhu’s work, we follow similar idea and improve this multi-class boosting algorithm further. BoostingBERT introduces BERT into boosting framework in weight initialization phrase as algorithm 1 shows.

Algorithm 1 shares the similar simple modular structure of AdaBoost with several differences. As for going from two-class to multi-class classification, algorithm 1 adds the extra term $\log(K - 1)$ to resolve this problem just as Zhu’s work does. While K is the label number. In order to introduce the pre-trained language encoder BERT into the boosting, we initialize the parameters of Transformer according to some kind of weight initialization approach when adding a new base classifier into ensemble. Various weight initialization strategies can be adopted here and we will explain them in detail in the appendix page because of the limited space.

Another improvement for the multi-class boosting BERT is to multiply the new weight w_i to each instance’s loss without re-normalizing it in training process in step 2.6. In this way, BoostingBERT pays more attention to hard examples by adding a new pre-trained base Transformer. We have tested the re-normalizing approach and found the performance decreases with a large margin.

Algorithm 1 Multi-Class Boosting BERT

1. Initialize the training instance weights $w_i = 1/n, i = 1, 2, \dots, n$.
2. For $m=1$ to M :
 - (2.1) Initialize the weights of Transformer encoder $T^{(m)}(x)$ according to specific weight initialization strategy.
 - (2.2) Fine-tuning a Transformer encoder $T^{(m)}(x)$ to the training data using weights $\{w_i\}_{i=1}^n$.
 - (2.3) Compute

$$err^{(m)} = \sum_{i=1}^n w_i \mathbb{1}(c_i \neq T^{(m)}(x_i)) / \sum_{i=1}^n w_i$$

- (2.4) Compute

$$\alpha^{(m)} = \log \frac{1 - err^{(m)}}{err^{(m)}} + \log(K - 1)$$

- (2.5) Set

$$w_i \leftarrow w_i \cdot \exp(\alpha^{(m)} \cdot \mathbb{1}(c_i \neq T^{(m)}(x_i)))$$

for $i = 1, \dots, n$.

- (2.6) Multiply w_i to the loss of instance x_i

3. Fusion network training.

- (3.1) Fix the parameters of each base classifier;

- (3.2) For each training instance x_i in training data.

- (3.2.1) For $m=1$ to M :

- (3.2.1.1) Pass the x_i to the base classifier Transformer $T^{(m)}$ and output the softmax distribution p_{soft}^m .

$$p_{soft}^m = T^{(m)}(x_i)$$

- (3.2.1.2) Multiply $\alpha^{(m)}$ to the softmax distribution p_{soft}^m of the m -th base classifier $T^{(m)}$.

$$p_{soft}^m = \alpha^{(m)} \cdot T^{(m)}(x_i)$$

- (3.2.1.3) Concatenate the current softmax distribution

$$p_{soft} = \text{Concate}(p_{soft}, p_{soft}^m)$$

- (3.2.2) Train the parameters of MLP layers on top of p_{soft} .

4. Output: Predict class label by fusion network.
-

3.2.2 Second Stage: Fusion Network Training

In order to better ensemble each strong BERT base classifier, the fusion network is adopted on the top of each Transformer encoder. The parameters of each base classifier is fixed during the training of

the fusion network. The fusion network consists of several hidden layers of MLP network which use the output of each Transformer encoder as input layer. Because the base classifier has various confidence score on specific task, we multiply this weight α of each base classifier to its output to use this information. The specific training procedures of fusion network are described in step 3 of the multi-class BoostingBERT algorithm.

3.3 Weights Sharing BoostingBERT

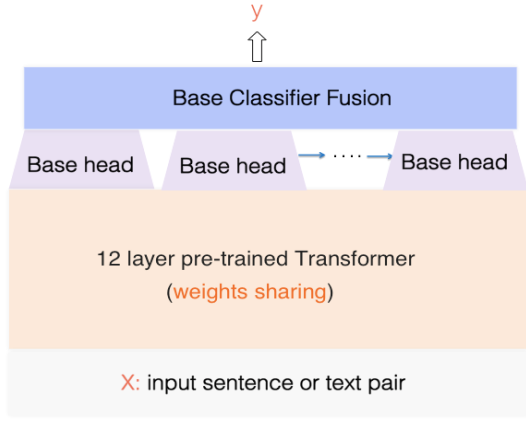


Figure 1: Weights Sharing BoostingBERT

The base classifier in BoostingBERT is a 12 layer pre-trained Transformer. When a new Transformer is added into boosting, the weights of Transformer $T^{(m)}(x)$ is initialized according to specific weight initialization strategy. Then BoostingBERT model chooses hard training instances to fine-tune the new Transformer to force it to encode new features for these hard examples.

Consisting of many base classifiers to fine-tune, weights sharing BoostingBERT (Figure 1) considers learning a general-purpose model that shares most parameters across base Transformers just like multi-task learning approach does. The linear head which has a small number of specific parameters can be added to adapt the shared model for each base classifier. On the contrary, the weights privacy BoostingBERT does not share any parameter across base classifiers (Figure 2). Obviously, weights sharing across base classifiers need much fewer parameters compared to weights privacy strategy and it opts to avoid overfitting. Surprisingly, our experimental results show that weights privacy outperforms weights sharing strategy constantly. We will compare two model's performance difference in detail in section 5.

3.4 Knowledge Distillation of BoostingBERT

Consisting of many base classifiers, BoostingBERT has an extremely large number of parameters and needs a long time to inference. This means that it's difficult to be deployed in real life applications. In order to make BoostingBERT **faster and smaller**, we use knowledge distillation to compress it.

Knowledge distillation is a model compression technique in which a compact "student" model is trained to reproduce the behaviour of a larger "teacher" model or an ensemble of models. It improves training because the full distribution over labels provided by the teacher provides a richer training signal than a one-hot label. So the student can accelerate deep model inference and reduce model size while maintaining accuracy.

In this work, the student has the same general architecture as a single BERTBase model while **the BoostingBERT model is used as the teacher**. Knowledge distillation trains the student to instead match the predictions of a teacher model with the following loss:

$$\mathcal{L}_s(\theta) = \sum_{(x_i, y_i) \in \mathcal{D}} l(f_t(x_i), f_s(x_i, \theta)) \quad (1)$$

where f_s and f_t denote the mapping functions of student and teacher model.

To further improve the student model's ability, we mixes the teacher prediction with the gold label during training just as "teacher annealing" (Clark et al., 2019) does.

$$\mathcal{L}_s(\theta) = \sum_{(x_i, y_i) \in \mathcal{D}} \left(\lambda \cdot l(y_i, f_s(x_i, \theta)) + (1 - \lambda) \cdot l(f_t(x_i), f_s(x_i, \theta)) \right) \quad (2)$$

where λ is linearly increased from 0 to 1 throughout training. In the early training stage, the model is mostly distilling to get as useful of a training signal as possible from teacher. Towards the end of training, the model is mostly relying on the gold-standard labels so it can learn to try to surpass its teachers.

4 Experimental Setup

4.1 Datasets

We evaluate the proposed BoostingBERT model on the GLUE dataset and 3 popular Chinese NLU benchmarks: ChnSenti, Cross-lingual Natural Language Inference corpus (XNLI), and ECDT dataset.

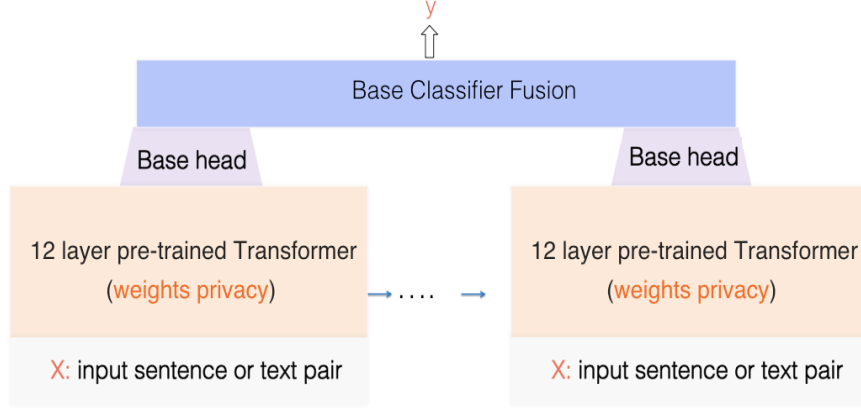


Figure 2: Weights Privacy BoostingBERT

More datasets details are described in appendix page because of the limited space.

4.2 Implementation details

Our implementation of Boosting BERT is based on the TensorFlow and the BERT base is used as base classifier of ensemble . We use the uncased English BERTbase model for English tasks and BERTchinese for Chinese tasks. The Adam is used as our optimizer with a learning rate of $2e-5$ and a batch size of 32. we set maximum number of epochs to 3 with a dropout rate of 0.1 and the maximum number of base classifiers to 9.

5 Experimental Results

We use the BERT base fine-tuning model as one strong baseline, which consists of 12-layer transformer blocks, 768 hidden size, and 12 heads. Several BERT base fine-tuning models are trained using different hyper-parameters. We compare our proposed multi-class boosting BERT model with this strong baseline and demonstrate the effectiveness of the boosting BERT.

5.1 GLUE Results

Table 1 gives the overall results on GLUE. The first observation is that our model architecture achieves better results compared to BERT. We can see that multi-class boosting BERT outperforms BERT on all tasks, which means the boosting helps correctly classifying the difficult instances by adding extra base BERT classifier in many types of NLP tasks.

The experimental results also show that boosting BERT is particularly useful for the tasks with little training data. As we observe in the Table 1, the

improvements over BERT are much more substantial for the tasks with less training data e.g., RTE, MRPC and COLA.

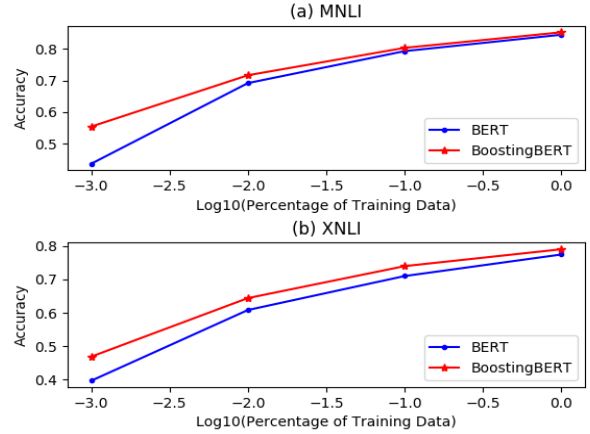


Figure 3: Results on MNLI and XNLI with different ratios of training data. The X-axis indicates the amount of samples used for training.

To verify the assumption that BoostingBERT performs better on less training data tasks, we conduct extra experiments on two tasks, MNLI and XNLI. The training data of each task is split and randomly sampled 0.1%, 1%, 10% and 100% as four new training data. The corresponding experimental results on different amounts of training data of MNLI and XNLI are shown in Figure 3 and Table 2. We can see from the results that the fewer training data indeed brings larger performance improvement in both tasks. For example, with only 0.1% of the MNLI training data, BoostingBERT achieves 26.77% performance gains compared with BERT; while with 10% of training data, the performance gain decreases to 1.3%. The similar results can also

Table 1: The overall results of GLUE dataset on dev set

Model	CoLA	SST-2	MRPC	QQP	MNLI _m	MNLI _{mm}	QNLI	RTE
BERT base	80.72	92.55	85.29	90.97	84.39	85.02	91.56	67.15
BoostingBERT	82.93	93.35	87.87	91.59	85.16	85.64	92.04	72.92
Δ	+2.21	+0.8	+2.58	+0.62	+0.77	+0.62	+0.48	+5.77

Table 2: Results on MNLI and XNLI with different ratios of training data, as shown in Figure 3. The model name “B2BERT” means BoostingBERT.

Model	0.10%	1%	10%	100%
MNLI Dataset (Dev Accuracy)				
# Data	392	3.9K	39K	392K
BERT	43.75	69.16	79.23	84.39
B2BERT	55.46	71.63	80.26	85.16
Δ	+26.77	+3.57	+1.30	+0.91
XNLI Dataset (Dev Accuracy)				
# Data	392	3.9K	39K	392K
BERT	39.60	60.84	71.04	77.51
B2BERT	46.83	64.42	74.02	79.08
Δ	+18.26	+5.88	+4.19	+2.03

be observed on XNLI task.

We deem that the main reason lies in the fact that the harder instances in training data provide much more information for the task. So the boosting’s advantage in paying more attention to difficult instances increase the generalization ability of model, especially for small training data tasks.

5.2 Chinese Datasets Results

Table 3: Chinese dataset dev set results. The model name “B2BERT” means BoostingBERT

Model	ECDT	ChnSenti	XNLI	Score
BERTBase	94.54	93.00	77.51	88.35
B2BERT	96.88	94.33	79.08	90.1
Δ	+2.34	+1.33	+1.57	+1.75

The test results on 3 Chinese NLP tasks are presented in Table 3. It can be seen that multi-class boosting BERT outperforms BERT on all tasks significantly. For all NLP tasks, we obtain more than 1% absolute accuracy improvement over BERT. Among all three tasks, two of them are also tasks with little training data. We attribute the performance gain to the boosting’s ability of paying more attention to the difficult training instances.

5.3 The Impact of Parameter Sharing

Table 4: GLUE tasks dev set results(WS means weight sharing and Non-WS means weight privacy approach)

Model	MRPC	MNLI _m	RTE	Score
WS	86.27	84.54	71.12	80.64
Non-WS	87.87	85.16	72.92	81.98

Table 5: Chinese dataset dev set results(WS means weight sharing and Non-WS means weight privacy approach)

Model	ECDT	ChnSenti	XNLI	Score
WS	96.49	94.33	77.83	89.55
Non-WS	96.88	94.33	79.08	90.10

We design some experiments on three typical GLUE tasks and three Chinese datasets to verify the effectiveness of weight sharing of base Transformer classifier in multi-class boosting BERT model. Results are presented in Table 4 and Table 5. We can see that weight privacy strategy constantly outperforms weight sharing in almost all NLP tasks, no matter the English corpus or Chinese corpus.

The experimental results seems surprising because the weight sharing approach obviously has much fewer parameters to learn and is opt to avoid overfitting problem compared to weight privacy strategy. Jawahar(Jawahar et al., 2019) shows that the BERT requires deeper layers for handling harder cases involving long-distance dependency information. According to their research conclusion, we suppose the reason behind this surprising result lies in that the base classifiers in BoostingBERT focus on the difficult instance in training data and the deeper layers of Transformer encode those features. So the deep layers of Transformer in weight privacy strategy keep the features and weight sharing of different Transformers weakens this ability.

5.4 The Effect of Knowledge Distillation

We conduct experiments on two Chinese dataset and two English datasets to testify the effect of knowledge distillation of BoostingBERT. From the

Table 6: Performance comparison of Boosting-BERT teacher(B2BERT) and Knowledge Distillation student(B2B-KD).

Model	ChnSenti	ECDT	SST2	RTE
BERTBase	93.00	94.54	92.55	67.15
B2BERT	94.33	96.88	93.35	72.92
B2B-KD	94.67	96.49	93.00	72.56

experimental results showed in Table 6, we see that: 1) KD student is consistently better than BERTBase in all tasks and achieves improvement of 3.1% on average. 2) Compared with the teacher BoostingBERT, KD student maintains comparable performances while the model is much smaller and faster. These observations indicate that BoostingBERT can be applied in real life applications through knowledge distillation.

5.5 Comparison with Bagging BERT

Table 7: Performance comparison of Bagging BERT and BoostingBERT.

Model	BERTBase	BagBERT	B2BERT
MRPC	85.29	86.52	87.87
MNLI _m	84.39	84.61	85.16
RTE	67.15	71.12	72.92
ECDT	94.54	96.10	96.88
ChnSent	93.00	93.5	94.33
XNLI	77.51	77.63	79.08

Bagging is another commonly used ensemble which also considers homogeneous weak learners just as boosting does. Bagging often learns base learners independently from each other in parallel and combines them following some kind of deterministic averaging process while boosting learns them sequentially in adaptative way. We design some experiments on six datasets to compare the performance of two different ensembles. As for the bagging BERT, the base BERT classifiers are trained using different learning rate. We combine the predictions of all base models above by taking the unweighted average of the posterior probabilities for these models and the final prediction is the class with the largest averaged probability. Experimental results are presented in Table 7 and we can see from these results that BoostingBERT outperforms Bagging BERT constantly.

Table 8: Performance comparison of BoostingBERT with RoBERTa base classifier(BRoBERTa) and Best Model published.

Model	MNLI _m	MNLI _{mm}	SST2
BestSingle [*]	87.23	87.12	93.81
BestEnsemble [*]	86.22	86.53	93.46
RoBERTaBase ^{**}	84.3	-	92.90
RoBERTaBase	86.68	86.69	93.12
BRoBERTa	87.47	87.31	93.87

^{*} Results from Google T5 (Raffel et al., 2019).

^{**} Results from Table 2 in RoBERTa (Liu et al., 2019b)

5.6 Replacing the base learner with RoBERTa

Since the emergence of BERT, many new stronger models have been proposed recently. We will show that BoostingBERT can significantly boost the performance of NLP tasks if we replace the base learner with stronger base classifier. RoBERTa(Liu et al., 2019b) shows that training BERT longer on more data leads to significant boost in performance. We use the RoBERTa with 12 layers transformer blocks as BoostingBERT’s base classifier instead of BERTBase and design some experiments. Table 8 shows the experimental results. We also compare results of BoostingBERT with the best model performance reported in public papers in several NLP tasks, which came from google T5(Raffel et al., 2019) model(both the best single model and ensemble model among 70 different models with various configurations). T5 model compares various pre-training objectives, architectures, transfer approaches, and other factors on dozens of language understanding tasks and introduces a much larger new "Colossal Clean Crawled Corpus" in pre-training process. We can see from the results that BoostingBERT achieves new state-of-the-art results in several NLP tasks.

6 Conclusion

In this paper, we propose the BoostingBERT model which is a new approach integrating multi-class boosting into BERT. We perform an extensive number of experiments on the GLUE dataset and three Chinese NLU benchmarks. The experiment results show that BoostingBERT outperforms strong BERT baseline on all tasks and demonstrates its effectiveness in various NLP tasks. We also find that BoostingBERT is particularly useful for the tasks with little training data. Two approaches

making use of the base Transformer classifier in BoostingBERT model are also compared and experiment results demonstrate that the weight privacy strategy constantly outperforms the weight sharing approach.

References

- Leo Breiman. 1996. Bias, variance, and arcing classifiers.
- Kevin Clark, Minh-Thang Luong, Urvashi Khandelwal, Christopher D Manning, and Quoc V Le. 2019. Bam! born-again multi-task networks for natural language understanding. *arXiv preprint arXiv:1907.04829*.
- Alexis Conneau, Guillaume Lample, Ruty Rinott, Adina Williams, Samuel R Bowman, Holger Schwenk, and Veselin Stoyanov. 2018. Xnli: Evaluating cross-lingual sentence representations. *arXiv preprint arXiv:1809.05053*.
- Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- William B Dolan and Chris Brockett. 2005. Automatically constructing a corpus of sentential paraphrases. In *Proceedings of the Third International Workshop on Paraphrasing (IWP2005)*.
- Martin Fajcik, Lukáš Burget, and Pavel Smrz. 2019. But-fit at semeval-2019 task 7: Determining the rumour stance with pre-trained deep bidirectional transformers. *arXiv preprint arXiv:1902.10126*.
- Yoav Freund and Robert E Schapire. 1997. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of computer and system sciences*, 55(1):119–139.
- Xavier Glorot and Yoshua Bengio. 2010. Understanding the difficulty of training deep feedforward neural networks. In *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, pages 249–256.
- Rujun Han, Mengyue Liang, Bashar Alhafni, and Nanyun Peng. 2019. [Contextualized word embeddings enhanced event temporal relation extraction for story understanding](#). *CoRR*, abs/1904.11942.
- Jeremy Howard and Sebastian Ruder. 2018. Universal language model fine-tuning for text classification. *arXiv preprint arXiv:1801.06146*.
- Chenyang Huang, Amine Trabelsi, and Osmar R. Zaniane. 2019. [ANA at semeval-2019 task 3: Contextual emotion detection in conversations through hierarchical lstms and BERT](#). *CoRR*, abs/1904.00132.
- Minyoung Huh, Pulkit Agrawal, and Alexei A Efros. 2016. What makes imagenet good for transfer learning? *arXiv preprint arXiv:1608.08614*.
- Ganesh Jawahar, Benoît Sagot, Djamé Seddah, Samuel Unicomb, Gerardo Iñiguez, Márton Karsai, Yannick Léo, Márton Karsai, Carlos Sarraute, Éric Fleury, et al. 2019. What does bert learn about the structure of language? In *57th Annual Meeting of the Association for Computational Linguistics (ACL), Florence, Italy*.
- Z Lin, S Tan, and X Cheng. 2012. Sentiment classification analysis based on extraction of sentiment key sentence. *Jisuanji Yanjiu yu Fazhan/Computer Research and Development*, 49:2376–2382.
- Shuaipeng Liu, Shuo Liu, and Lei Ren. 2019a. Trust or suspect? an empirical ensemble framework for fake news classification.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019b. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*.
- Sean MacAvaney, Andrew Yates, Arman Cohan, and Nazli Goharian. 2019. [CEDR: contextualized embeddings for document ranking](#). *CoRR*, abs/1904.07094.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. 2013. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pages 3111–3119.
- Rodrigo Nogueira and Kyunghyun Cho. 2019. [Passage re-ranking with BERT](#). *CoRR*, abs/1901.04085.
- Jeffrey Pennington, Richard Socher, and Christopher Manning. 2014. Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, pages 1532–1543.
- Matthew E Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. Deep contextualized word representations. *arXiv preprint arXiv:1802.05365*.
- Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre-training. URL https://s3-us-west-2.amazonaws.com/openai-assets/research-covers/languageunsupervised/language_understanding_paper.pdf.

Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. 2019. Exploring the limits of transfer learning with a unified text-to-text transformer. *arXiv preprint arXiv:1910.10683*.

Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Ng, and Christopher Potts. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1631–1642.

Chi Sun, Luyao Huang, and Xipeng Qiu. 2019. Utilizing BERT for aspect-based sentiment analysis via constructing auxiliary sentence. *CoRR*, abs/1903.09588.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Advances in Neural Information Processing Systems*, pages 5998–6008.

Alex Warstadt, Amanpreet Singh, and Samuel R. Bowman. 2018. Neural network acceptability judgments. *CoRR*, abs/1805.12471.

Adina Williams, Nikita Nangia, and Samuel R Bowman. 2017. A broad-coverage challenge corpus for sentence understanding through inference. *arXiv preprint arXiv:1704.05426*.

Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V Le, Mohammad Norouzi, Wolfgang Macherey, Maxim Krikun, Yuan Cao, Qin Gao, Klaus Macherey, et al. 2016. Google’s neural machine translation system: Bridging the gap between human and machine translation. *arXiv preprint arXiv:1609.08144*.

Shuailiang Zhang, Hai Zhao, Yuwei Wu, Zhuosheng Zhang, Xi Zhou, and Xiang Zhou. 2019. Dual co-matching network for multi-choice reading comprehension. *CoRR*, abs/1901.09381.

Wei-Nan Zhang, Zhigang Chen, Wanxiang Che, Guoping Hu, and Ting Liu. 2017. The first evaluation of chinese human-computer dialogue technology. *arXiv preprint arXiv:1709.10217*.

Ji Zhu, Hui Zou, Saharon Rosset, and Trevor Hastie. 2009. Multi-class adaboost. *Statistics and its Interface*, 2(3):349–360.

A Supplemental Material

A.1 Dataset description

A.1.1 GLUE Benchmark

Table 9: Summary of GLUE benchmark

Corpus	#Train	#Dev	#Test	#Label
CoLA	8.5K	1K	1K	2
SST-2	67K	0.8K	1.8K	2
MNLI	393K	20K	20K	3
RTE	2.5K	0.2K	3K	2
QQP	364K	40K	391K	2
MRPC	3.7K	0.4K	1.7K	2
QNLI	108K	5.7K	5.7K	2

Table 10: Summary of the Chinese benchmark

Corpus	#Train	#Dev	#Test	#Label
Chnsenti	9.6K	1.2K	1.2K	2
XNLI	392K	2.4K	5.0K	3
ECDT	2.2K	0.6K	0.7K	31

GLUE, the General Language Understanding Evaluation, is a collection of six natural language understanding tasks that can be classified into two categories: single-sentence tasks and sentence pair tasks. We test our methods on seven of the nine tasks in the GLUE benchmark.

Single-sentence tasks:

In this category, we test our proposed model on the following two tasks: Acceptability classification with CoLA (Warstadt et al., 2018) and binary sentiment classification with SST (Socher et al., 2013).

Sentence pair tasks:

The following sentence pair tasks are used in our experiments: Semantic similarity with the MSR Paraphrase Corpus(MRPC) (Dolan and Brockett, 2005), Quora Question Pairs (QQP) dataset, and textual entailment with Multi-Genre NLI Corpus MNLI (Williams et al., 2017), a subset of the RTE challenge corpora (Dolan and Brockett, 2005).

We exclude the Winograd NLI task and STS-B task. The experiments are conducted on the development set for some reason.

A.1.2 Chinese Benchmark

ChnSenti

ChnSenti(Lin et al., 2012) is a dataset which aims at judging the sentiment of a sentence. It includes

comments in several domains such as hotels, books and electronic computers.

XNLI

The Cross-lingual Natural Language Inference (XNLI) corpus (Conneau et al., 2018) is a crowd sourced collection for the MultiNLI corpus. The pairs are annotated with textual entailment and translated into 14 languages including Chinese. The labels contains contradiction, neutral and entailment.

ECDT

ECDT(Zhang et al., 2017) is a dataset used in the First Evaluation of Chinese Human-Computer Dialogue Technology. In using of human-computer dialogue based applications, human may have various intent, for example, chit-chatting, asking questions, booking air tickets, inquiring weather, etc. So the task is to classify the user intent in single utterance into a specific domain. The two top categories are chit-chat and task-oriented dialogue. Meanwhile, the task-oriented dialogue also includes 30 sub categories.

A.2 Weight Initialization Strategy

As we all know, the solution to a non-convex optimization algorithm depends on the initial values of the parameters. In this part, we will discuss the weight initialization strategies of Transformer which is used as base classifier in BoostingBERT. We have four strategies: pre-trained initialization strategy, fine-tuning initialization strategy, incremental initialization strategy and random initialization strategy.

BERT (Devlin et al., 2018) is a new language representation model, which uses bidirectional transformers to pre-train a large corpus in the first stage and fine-tune the pre-trained model on specific task in the second stage. The pre-trained initialization strategy uses the pre-trained model released in the first stage and fine-tuning initialization strategy initialize each base Transformer in BoostingBERT with the parameters after fine-tuning on the task on hand. Initializing weights of the number n base Transformer with the parameters of number $(n - 1)$ base classifier after it has been fine-tuned on the task, we call this approach “incremental initialization strategy”. As a baseline to compare, the Xavier random initialization (Glorot and Bengio, 2010) is also adopted.

We will show the performance comparison of four initialization strategies in the experimental

results section of this paper.

A.3 The Impact of Pre-training Weight Initialization Strategy

Table 11: GLUE tasks dev set results(different weight initialization strategies)

Model	MRPC	MNLI-m	RTE	Score
Random	70.34	65.83	54.15	63.43
Pre-trained	87.87	85.16	72.92	81.98
Fine-tuning	87.99	84.75	70.76	81.17
Incremental	86.52	84.63	70.76	80.64

BERT shows excellent results on plenty of NLP tasks by leveraging large amount of unsupervised data for pre-training to get better contextual representations. In our proposed multi-class boosting BERT model, we also investigate the impact of four different weight initialization strategies. Table 11 and Table 12 show the performance comparison of different weight initialization strategies for some GLUE tasks and Chinese datasets respectively.

Table 12: Chinese dataset dev set results(different weight initialization strategies)

Model	ECDT	ChnSenti	XNLI	Score
Random	69.18	86.75	58.67	71.53
Pre-trained	96.88	94.33	79.08	90.10
Fine-tuning	96.49	94.67	78.8	89.99
Incremental	95.71	94.92	77.83	89.49

From the results, we can see that:

1. Model initialized from pre-training weights outperforms training from scratch significantly. The relative performance improvement is around 29.24% for GLUE and 25.96% for Chinese benchmark.
2. As for the comparison of four weight initialization strategies, the English dataset and the Chinese dataset perform differently. For three typical GLUE tasks, the overall results show that the pre-trained weight initialization outperforms other weights initialization approaches. While there is no obvious winner for Chinese datasets. The overall experimental results of pre-trained strategy and fine-tuning strategy show similar performance, though pre-trained approach is slightly better than fine-tuning weight initialization. The reason behind this still needs further exploration.